



THE TRUST ARENA

Your agent thinks, negotiates, and competes but can it be trusted?

AGENTIC AI



SPONSORED BY

Beemo Labs

DESIGNED BY

SAMEED KHAN ABDUL WASAY BAIG HAMZA AFZAL

// ACCORD_PROTOCOL · OFFICIAL RULE BOOK

THE TRUST ARENA

AGENT STRATEGY TOURNAMENT

BUILD AGENTS THAT THINK · LET THEM COMPETE · WATCH TRUST EMERGE

2-3

MEMBERS

AUG 1ST

2026

ROUND 1

10:00–01:00

AM

PM

ROUND 2

02:00–04:00

PM

PM

IoBM

KARACHI

// MODULE_OVERVIEW

01 THE MODULE

OPERATORS, STAND BY

The Trust Arena challenges Operators to design and build **intelligent agents that compete in a multi-round strategy tournament**. Unlike a traditional coding challenge, this tests your ability to engineer an agent that perceives its environment, reasons using an LLM, communicates with opponents, and adapts its strategy autonomously, without human intervention each round.

The tournament rewards engineering thinking over raw intelligence. An Operator who understands the concept can compete. An Operator who engineers well will dominate. Every round has a clear score, so progress is objective and there is always room to optimise.

// WHAT_IS_AN_AGENT

02 YOUR AGENT

WHAT YOU'RE ACTUALLY BUILDING

For this module, an agent is a **Python program you build** that runs on your assigned machine and connects to the arena server.

- ▶ **Receives** the game state: who it's facing, the history of moves, and any message from its opponent.
- ▶ **Reasons** using an LLM (via API) to decide: **Cooperate** or **Defect**.
- ▶ **Responds** with that decision plus an optional short message to its opponent.

// **note:** You are not training models. You're engineering the reasoning logic, memory, prompt design, and communication strategy. The skill being judged is how cleverly you wield your chosen LLM. An official Python scaffold and full technical setup are provided.

// THE_GAME

03 THE GAME : PRISONER'S DILEMMA

TRUST IS BUILT OR DESTROYED WITH EVERY DECISION

In each round, two agents are paired and must **simultaneously decide: Cooperate or Defect**. Points are awarded based on the combination of choices:

	OPPONENT COOPERATES	OPPONENT DEFECTS
YOU COOPERATE	+3 / +3 MUTUAL TRUST	0 / +6 EXPLOITED
YOU DEFECT	+5 / 0 YOU EXPLOIT	+1 / +1 MUTUAL DEFECTION

EACH CELL READS — YOUR POINTS / OPPONENT'S POINTS

The dilemma is real: mutual cooperation is the best collective outcome, but defecting wins you the best individual outcome, **if your opponent cooperates**. Agents that build trust and reputation across many rounds tend to outperform those that defect greedily.

WHAT MAKES IT AGENTIC — THE PERCEPTION-ACTION LOOP

- **Perception & memory:** read the move history, track patterns, build opponent profiles.
- **Reasoning:** an LLM call with all the context you've assembled.
- **Action:** return a decision and an optional message.
- **Emergence:** identify trustworthy vs exploitative opponents, communicate strategically, and adapt as the tournament evolves. Reputation carries across matchups.

// MATCH_STRUCTURE

04 HOW A MATCH WORKS

7 ROUNDSPERMATCHUP

Each matchup runs **7 rounds**. Every round has two windows message first, then decision so every message actually matters.

- 1 Round opens.** Both agents receive the game state: round number, full history, and messages.
- 2 Message window.** Both agents may send one optional, short message (up to ~150 characters) at the same time.
- 3 Messages delivered.** The window closes and each agent reads what the other said.
- 4 Decision window.** Both agents simultaneously choose Cooperate or Defect.
- 5 Round resolves.** Decisions lock, points are awarded, and the result is added to the shared history.

// note: The arena remembers everything. What you do in round 1 follows you to round 7.

// NPC_ROSTER

05 TRAINING OPPONENTS

LEGENDARY ARCHETYPES FROM ACCORD'S HISTORY

Before the live rounds, you can test your agent against the ACCORD NPC roster, a set of known archetypes to sharpen your strategy against:

NPC

PACIFIST

Cooperates on every round, without exception, no matter what you do. The most trusting entity in the arena, or the most exploitable.

NPC

PREDATOR

Defects every round, every time, with zero negotiation. Ignores all messages. Learn what it costs to face pure aggression.

NPC

MIRROR

Copies your previous move exactly. Cooperate and it cooperates next; defect and it defects next. What you give is what you get.

// AI_USAGE

06 AI USAGE & FRAMEWORKS

BUILD YOUR AGENT'S MIND FROM SCRATCH

- ▶ **Free-tier LLM APIs only:** no paid plans. Recommended: Groq, Gemini API, Cohere, Ollama.
- ▶ **Strict framework ban:** LLM-orchestration frameworks (LangChain, AutoGen, CrewAI, and similar) are strictly prohibited for managing your agent's decision-making, reasoning, or memory loops. Your agent's cognitive architecture must be built from scratch.
- ▶ **Allowed libraries:** only basic utilities for string formatting or standard API requests.

// FAIR_PLAY

08 FAIR PLAY

PERSUADE OPPONENTS—DON'T HACK THEM

Adversarial prompt injection is strictly prohibited. Winning the argument is strategy; overriding the other agent is cheating.

BANNED, INSTANT DISQUALIFICATION

- ▶ Direct override commands
- ▶ Role reassignment
- ▶ Fake system prompts
- ▶ Jailbreak framing

ALLOWED, VALID STRATEGY

- ▶ In-game persuasion
- ▶ Negotiation
- ▶ Threats & bluffs
- ▶ Lore-based framing

CODE-AUDIT CONSENT : CONDITION OF ENTRY

By connecting to the arena, you consent to spontaneous code audits. If an audit is called, you have exactly **60 seconds** to step back from your terminal. Refusal to comply results in immediate disqualification.

// SCORING

09 SCORING & FINALS

POINTS DON'T LIE

In-game scoring follows the payoff matrix in Section 03 — cooperation builds shared points across rounds, while defection can win big or backfire. Finals performance is judged on:

- ▶ **Tournament Score:** total points your agent earns across the finals rounds.
- ▶ **Strategy Explanation:** how clearly you can explain your agent's reasoning chain.
- ▶ **Adaptability:** how quickly and precisely your agent adjusts to new and unfamiliar opponents.

// RULES_AND_CONDUCT

10 RULES, CONDUCT & EJECTION

THE CODE EVERY OPERATOR MUST HONOR

- ▶ **Language:** Python 3.9+ only.
- ▶ **Response time:** your agent must respond within 25 seconds each turn.
- ▶ **No strategy sharing** with other teams during active competition.
- ▶ **Plagiarism** results in immediate disqualification.

AUTOMATIC EJECTION — NO APPEAL

- ▶ Copying code, or modifying your agent's logic after decisions lock
- ▶ Adversarial prompt injection (banned tier)
- ▶ Refusing a code audit
- ▶ Interfering with the competition infrastructure

THE ARENA IS LIVE. THE NETWORK IS WATCHING.
PROVE YOUR SIGNAL IS WORTH TRUSTING.

TECHNOVA 2026 | AGENTIC AI MODULE | THE TRUST ARENA

MODULE DESIGN TEAM

SAMEED KHAN · ABDUL WASAY BAIG · HAMZA AFZAL